

Dahan, Amir Michael  
Supervised by Prof. Margarita Osadchy

Feb. 2022

M.Sc. Final Project Report - *Medical Face Mask Detection*

## Contents

<b>1 Goals</b>	<b>3</b>
1.1 Expected Achievements . . . . .	4
<b>2 Data</b>	<b>5</b>
2.1 Datasets . . . . .	5
2.1.1 Face Mask Detection Dataset [1] . . . . .	5
2.1.2 Face Mask Dataset [1] & Natural Images Dataset [2] . . . . .	5
2.2 Transformations used for data . . . . .	6
<b>3 Architecture</b>	<b>7</b>
3.1 The Model . . . . .	7
3.2 Training . . . . .	8
3.3 Testing Pipeline . . . . .	8
<b>4 Results</b>	<b>9</b>
<b>5 Gui</b>	<b>11</b>
<b>6 Bibliography</b>	<b>12</b>
<b>7 Appendixes</b>	<b>12</b>
7.1 The Code . . . . .	13
7.2 How To Run The Project . . . . .	13
7.3 Demo Agena . . . . .	14
7.3.1 Textual . . . . .	14
7.3.2 Gui . . . . .	15
7.3.3 Video . . . . .	16

## 1 Goals

In this project, the goal is to label all faces in the given image as mask/no mask

We'll need to determine which of these women is wearing a medical mask.



FIGURE 1: Original Image



FIGURE 2: Annotated Image

## 1.1 Expected Achievements

Contents:

1. Take both datasets and use augmentation to improve pictures.
2. Pick two models best suits (CNN) our problem & search for the best hyper parameters.
3. Train two models on datasets and save it for later use
  - Fine-tune the human detector model's upper layer
  - Train mask detector from scratch
4. Test it using a GUI or a great integration script (see appendixes section 7 ), and run on an unseen test images

## 2 Data

Data used for this project:

### 2.1 Datasets

#### 2.1.1 Face Mask Detection Dataset [1]

Comprised of the following three classes:

- Face with mask
- Face without mask
- Mask worn incorrectly

Classes Distribution: (Equal number of images per class, to eliminate bias)

- mask\_weared\_incorrect: 2997
- with\_mask: 2997
- without\_mask: 2997

#### 2.1.2 Face Mask Dataset [1] & Natural Images Dataset [2]

A custom superset of both [1] & [2], Comprised of the following classes:

- airplane
- car
- cat
- dog
- flower
- fruit
- mask\_weared\_incorrect
- motorbike

- with\_mask
- without\_mask

Classes Distribution:

- airplane: 727
- car: 968
- cat: 885
- dog: 702
- flower: 843
- fruit: 1000
- mask\_weared\_incorrect: 2997 -> 1000
- motorbike: 788
- with\_mask: 2997 -> 1000
- without\_mask: 2997 -> 1000

## 2.2 Transformations used for data

- Random Horizontal Flip: Horizontally flip the given image randomly with a given probability
- Random Resized Crop: Crop the given image to random size and aspect ratio

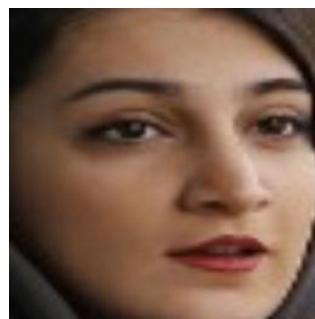


FIGURE 3: Cropping Sample

## 3 Architecture

In this section I'll describe the model and the modifications I've made during second phase.

### 3.1 The Model

I chose a CNN, no other choice for that kind of task. I started from a known network that I previously used for a similar task, combined with an intuition from a known model for image classification. I played a little with the number of filters for each layer, this way I can capture more features. The problem with this method was that it takes a lot of RAM (over 4GB), which made me debug on a pre-trained resnet18.

```
CNN(
    (loss_func): CrossEntropyLoss()
    (feature_extractor): Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): ReLU()
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU()
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Flatten(start_dim=1, end_dim=-1)
        (1): Linear(in_features=200704, out_features=1024, bias=True)
        (2): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (3): ReLU()
        (4): Linear(in_features=1024, out_features=512, bias=True)
        (5): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (6): ReLU()
        (7): Linear(in_features=512, out_features=10, bias=True)
    )
)
```

FIGURE 4: architecture

```

SGD (
    Parameter Group 0
        dampening: 0
        lr: 0.001
        momentum: 0.9
        nesterov: False
        weight_decay: 0
)

```

FIGURE 5: optimizer

### 3.2 Training

The training phase includes three models:

1. Natural Image:
  - A Pre-Trained Model that detects human beings (alongside other)
  - Isn't being used for final classification
  - Used for Human Model Retraining
2. Human Model
  - Fine tune 1 model's top layer
  - Classes: elaborated in 2.1.2
  - Add three classes (with\\_mask, without\\_mask, mask\\_weared\\_incorrect)
  - Remove 'person' class in return
  - Used to determine if the object is a person
3. Face Mask Detection
  - A model we train from scratch
  - Classes: elaborated in 2.1.1
  - Used for Human Model Retraining

### 3.3 Testing Pipeline

- Take an image
- Determine if the object/s in the picture are human (using Human Detection Model model)

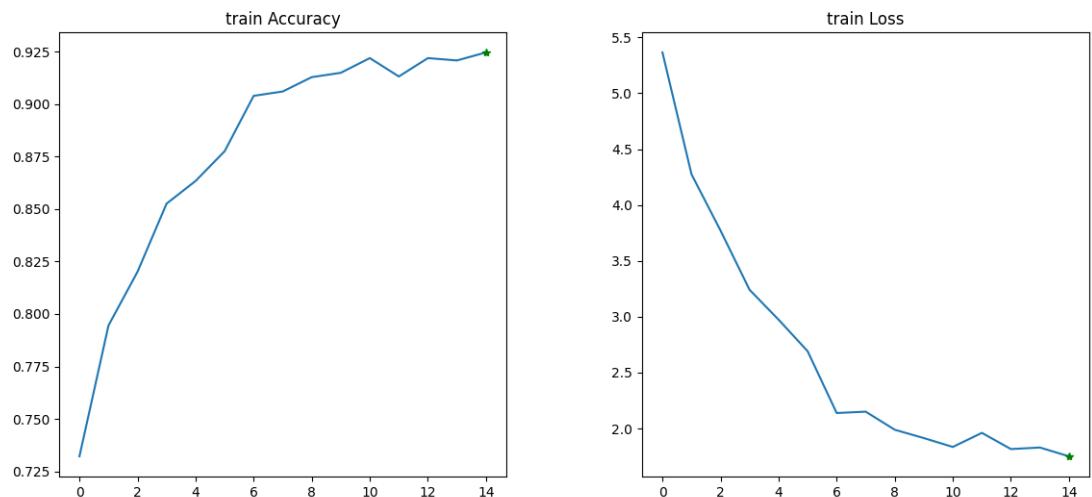
- Crop the object/s one by one (using first model), & determine masked/non-masked/partially-masked (using second model)

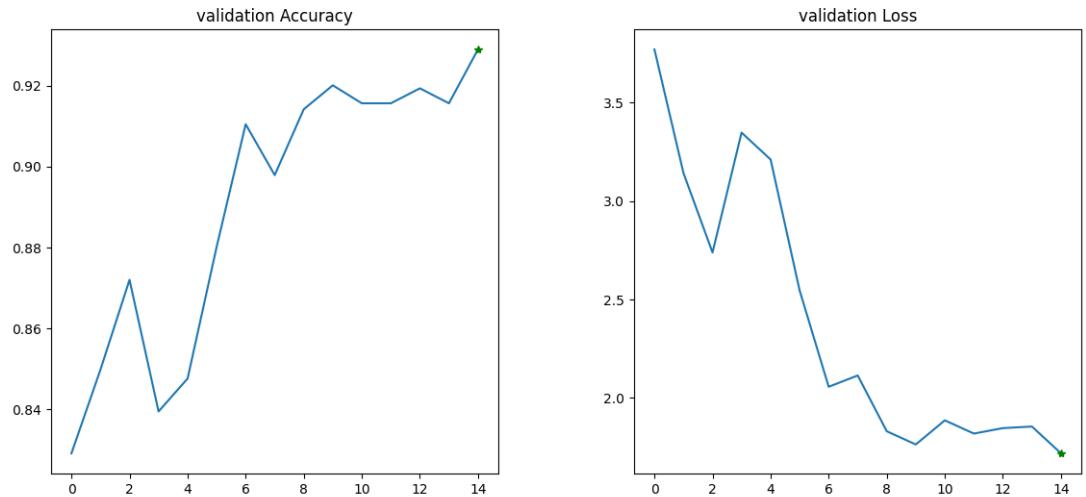
## 4 Results

Results will be elaborated for both models:

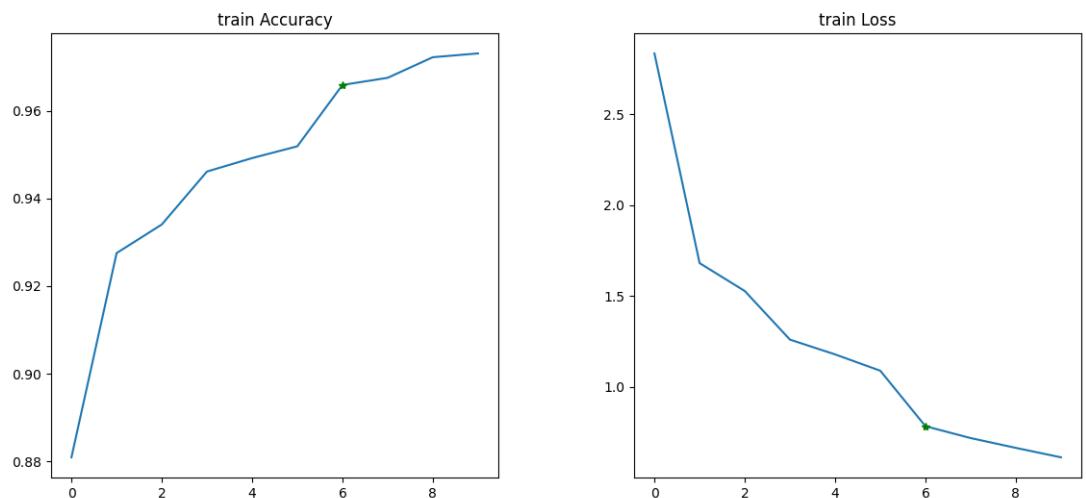
- Human Detection - Accuracy: 97.54%, Loss: 0.57
- Medical Face Mask Detection - Accuracy: 93.32% Loss: 1.71

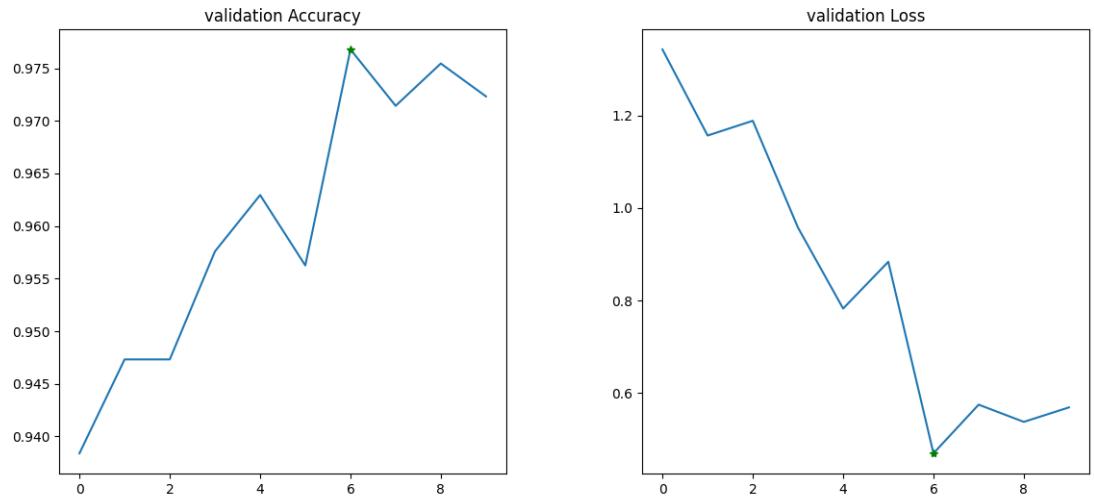
### 1. Mask Detection





## 2. Human Detection





## 5 Gui

In this section we'll describe how our gui looks like and some screenshots

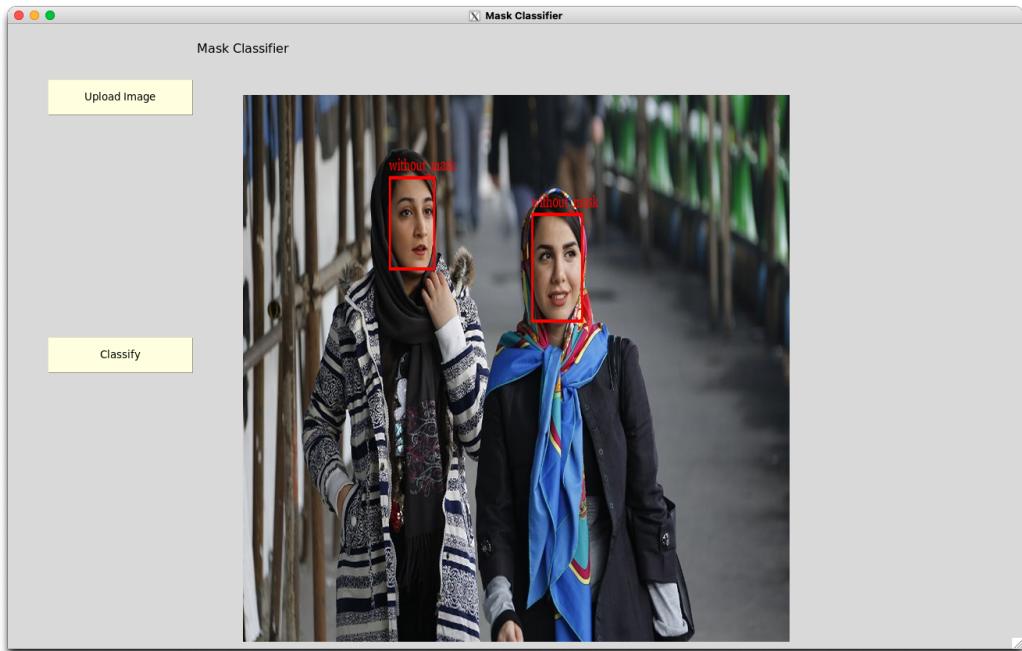


FIGURE 6: Non-Masked Sample

## 6 Bibliography

### References

- [1] Face Mask Detection, Building a face mask classifier  
<https://www.kaggle.com/vijaykumar1799/face-mask-detection>
- [2] A compiled dataset of 6899 images from 8 distinct classes.  
<https://www.kaggle.com/prasunroy/natural-images>
- [3] FaceNet: A Unified Embedding for Face Recognition and Clustering  
<https://arxiv.org/abs/1503.03832>

## 7 Appendixes

### 7.1 The Code

My project is consisted of the following components:

---

<b>doc</b>	Documentation, source and pdf
<b>out</b>	Saved Checkpoint - model & state dict.
<b>samples</b>	Test images for Demo
<b>src</b>	Project's sources
<b>support</b>	Dockerfile to reproduce the required env.

---

TABLE 1

---

<b>Camera.py</b>	Capture frames from Camera or a Video
<b>FaceMaskClassificationUtils.py</b>	Shared Helper and constants
<b>FaceMaskDetection.py</b>	FaceMask classifier model
<b>FinalProject.py</b>	train and classify without a GUI
<b>Gui.py</b>	A useful GUI assisting upload and classify images
<b>HumanDetection.py</b>	Human Detection Model
<b>ObjectCrop.py</b>	External Lib For Objects Detection

---

TABLE 2

### 7.2 How To Run The Project

- Train the model and test it, export a checkpoint file for the model and weights

```
## python src/FinalProject.py --help
optional arguments:
  -h, --help            show this help message and exit
  -H HUMAN_MODEL_PATH, --human_model_path HUMAN_MODEL_PATH
                        Human Model Path
  -a HUMAN_DATA_PATH, --human_data_path HUMAN_DATA_PATH
                        Human Data Path
  -M MASK_MODEL_PATH, --mask_model_path MASK_MODEL_PATH
```

```

        Face Mask Model path
-b MASK_DATA_PATH, --mask_data_path MASK_DATA_PATH
        Face Mask Data path
-F IMAGE_PATH, --image_path IMAGE_PATH
        image to classify
--train

```

- Run GUI to classify images selected by the user:

```

$$ python src/Gui.py --help
optional arguments:
-h, --help            show this help message and exit
-H HUMAN_MODEL_PATH, --human_model_path HUMAN_MODEL_PATH
        Human Model Path
-M MASK_MODEL_PATH, --mask_model_path MASK_MODEL_PATH
        Face Mask Model path

```

- Run Video/Camera capture and classify frames

```

$$ python src/Camera.py
usage: Camera.py [-h] -H HUMAN_MODEL_PATH -M MASK_MODEL_PATH -F VIDEO_PATH
                  [--skip_capturing]

```

Prerequisites:

- Pre trained model checkpoint (Model & weights)
- Project's Python source code

## 7.3 Demo Agena

### 7.3.1 Textual

```

$ cd ~/Project/ && time python src/FinalProject.py \
-H out_models/CombinedModel2.pth \
-M out_models/MaskModel2.pth \
-a data/human/ -b data/face-mask/ \

```

```
-N out_models/NaturalNode12.pth \
-c data/natural_images \
-F samples/nre/Partial.jpg
```

```
((venv36) adahan@kodkoda4:~/Project$ cd ~/Project/ && time python src/FinalProject.py -H out_models/CombinedModel2.pth -M out_models/MaskModel2.pth -a data/human/ -b data/face-mask/ -N out_models/NaturalNode12.pth -c data/natural_images -F samples/nre/Partial.jpg
Namespace(human_data_path='data/human/', human_model_path='out_models/CombinedModel2.pth', image_path='samples/nre/Partial.jpg',
mask_data_path='data/face-mask/', mask_model_path='out_models/MaskModel2.pth', natural_data_path='data/natural_images', natural_model_path='out_models/NaturalNode12.pth', train=False)
FinalProject: shouldn't train
human model successfully loaded
mask model successfully loaded
FinalProject: loading models took 0:00:02.425118
trying to classify image samples/nre/Partial.jpg
classify if it is human
cpu
classify_is_human: this is a with_mask.
classify_is_human: it is human
cropping image
single_image_classify got 1 faces to classify
classify_mask_usage: this is a mask_woread_incorrect.
FinalProject: testing image took 0:00:00.381641

real    0m6.211s
user    0m5.545s
sys     0m2.538s
-
```

FIGURE 7: Cropping Sample

### 7.3.2 Gui

```
$ cd ~/Project/ && python src/Gui.py \
-H out_models/CombinedModel2.pth \
-M out_models/MaskModel2.pth
```

I'll attach here samples of my self Gui tests:

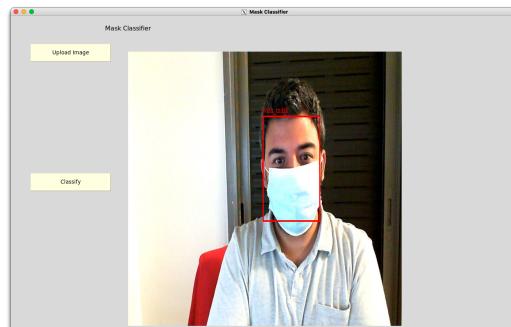


FIGURE 8: Amir is with his mask ON



FIGURE 9: Amir is with his mask Partially

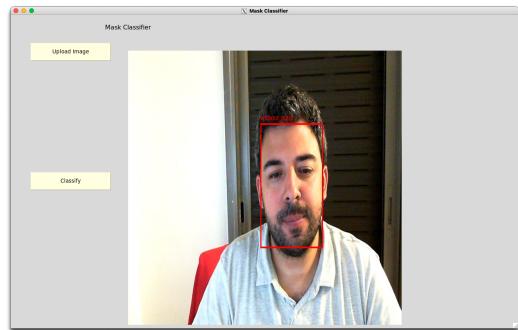


FIGURE 10: Cropping Sample

### 7.3.3 Video

```
$ cd ~/Project/ && python src/Camera.py \
    -H out_models/CombinedModel2.pth \
    -M out_models/MaskModel2.pth \
    -F samples/nre/Movie.mov --skip
```